Great post! The lack of hard numerical properties that be easily compared is an interesting observation.

I think the main takeaway is that, due to the vast number of different potential properties at this layer, it puts more focus onto the actual applications on the chain rather than the chain specs (e.g. people are more likely to choose a Mac because of GarageBand rather than because of the CPU specs).

Since application UX is mostly dependent on being cheap and fast, chains that instead prioritize decentralization, censorship, and privacy may have a harder time gaining adoption.